

Attendance \Rightarrow 101.

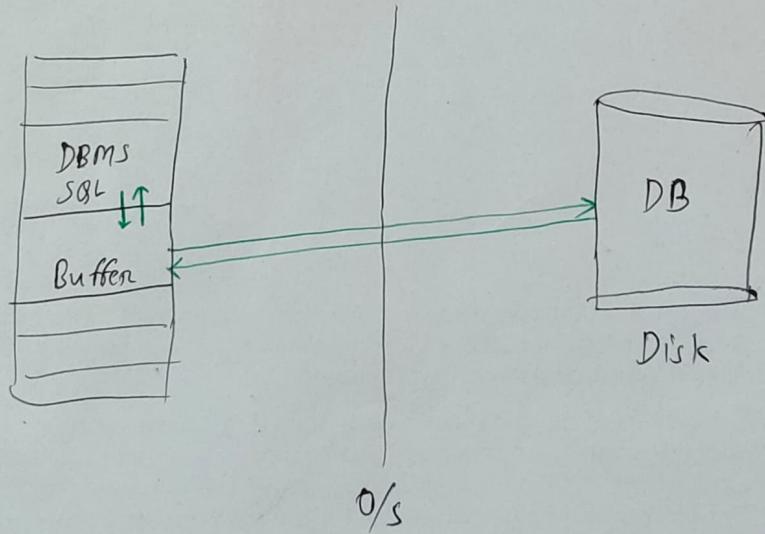
Quiz \Rightarrow 2 } Best 2 \Rightarrow 201.
Assignment = 1 }

Midterm \Rightarrow 301.

Final \Rightarrow 401.

Book \Rightarrow system concept 7th
edition.

④ Disk Based DBMS System



④ Data read/write through Buffer.

- SQL runs on Buffer.

④ SQL first search on Buffer, if not found, take it into buffer then provide output.

④ Disk access time in millisecond.

But memory access time in nano second.

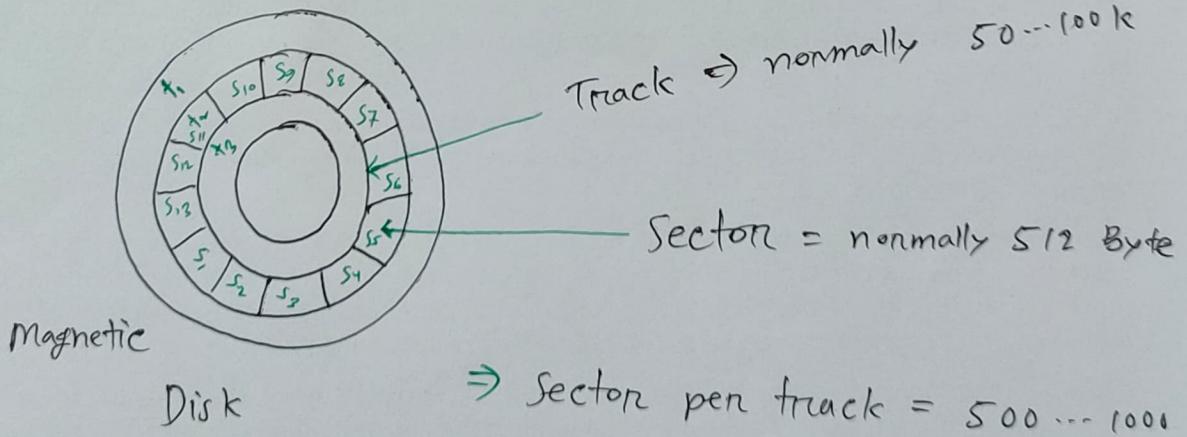
Storage architecture

⇒ two type of storage.

- Volatile ⇒ primary memory
 - content will be erased when power off.
- Non-Volatile ⇒ secondary memory
 - content will remain..

Storage Hierarchy - [Slide-12]

- Top 4 memory type are related to the database.
 - Direct addressable
- Bottom two are not possible to access directly or randomly. These need sequential access.
 - Sequential addressable.



- often track have more sector than innen track



Student (ID, Name, street, city)

Address : Track - 500 & sector - 200

Query : SELECT *
FROM Student

| Current head position in track - 1.
| No data in Buffer.

⇒ What will be the access time?

→ Access time = Start from "SQL issue" and end in
"Start of data transfer."

Access time = Head Reposition time from track 1 to 500

+

time to place the head at sector - 200
need to rotate the disk.

Track reposition time = seek time (4-10 milliseconds)

5-20 milliseconds
Sector reposition time = Rotational latency (4-11 milliseconds)

⑥ Data transfer rate: depends on hardware manufacturer.

⇒ 25 - 200 MBps
innen tracks outer tracks

⑦ Average seek time and rotational latency will be given, calculate in percentage.

If average given, then max is = $2 \times$ average.

As we know, average = $\frac{\text{best} + \cancel{\text{worst}}}{2} = \frac{0 + \text{max}}{2} \Rightarrow \text{max} = 2 \cdot \text{ave}$

RAID

⊗ Redundant Arrays of Independent Disks:

- use of multiple disk in parallel provides high capacity and high speed with high reliability.

⊗ Bit level striping:

Block Size
⇒ 4 kB, 8 kB,
16 kB.

- split the bits of each byte across multiple disks. one bit per disk
- seek/access time worse than for a single disk.
- Not so efficient for large data.

⊗ Block level striping:

i-block will go to $(id_{on}) + 1$

⊗ Raid level - 0:

- No redundant data
- only one copy of data, if a disk fail then data will be lose.
- But provide high performance.

RAID level - 1:

- Best write performance
- two copy of data, like a mirror disk

Given, disk size = 4 TB

How many disk will be needed to have storage 24 TB effective storage?

$$\text{RAID-0} \Rightarrow 6x$$

$$\text{RAID-1} \Rightarrow 12x (6+6)$$

RAID-5: Block striping distributed parity

- parity bit can help us to recover a data.

↳ odd parity
↳ even parity

- support one disk failure.

Recovery mechanism and updating mechanism.

L-03 / 02.02.2024/

Indexing

Student

ID	Name	Street	City
1	Abdullah		comilla
2	Abid		Dhaka
:	:		:
3000	Abid		Rajshahi
:	:		:
10000	Anif		Dhaka
:			:
20000	Abid		Dhaka

Query:

```
CREATE INDEX Name
ON Student (Name)
```

Search key	Pointers	Bucket
Abdullah	1	3000
Abid	multiple	20000
Anif	10,000	

Sorted

Syntax:

CREATE INDEX index-name
ON table-name (column1, column2...)

 Index file must be ordered or sorted / sorted.

 Two kinds of indices

- Ordered
- Hash

Index Evaluation Metrics

⇒ Insertion Time :

- if there are 50 attribute in a table then,
we need to insert the search key on 50
index first then we can insert in database.

⇒ Deletion Time :

- same as insertion

⇒ Access Time : depends on type.

point access ($=$)

range access ($<$, $>$)

 Space overhead for large table with so many
attribute.

L-04 / 04.09.2024

Indexing

Access type:

- point access $\Rightarrow (=)$
- range access $\Rightarrow (<,>)$

Primary Index / Clustering Index:

- orden of index file = orden of original data
- orden of index file \neq not necessarily the primary key.

Secondary Index:

- orden of index file \neq orden of original data

Dense Index:

one index/search key = one data

- index record appears for every search-key value in the file.

Query - 1:

```
SELECT *
FROM Instructor
WHERE ID = 83821
```

\Rightarrow Point access

Query - 2:

```
SELECT *
FROM Instructor
WHERE ID > 58583
```

\Rightarrow Range access

\Rightarrow Let one tuple/row/data = 1 Block

then, instructor relation size = 12 blocks.

Question - 1:

- Explain how Query-1 and Query-2 will be processed without index file?

Question - 2:

- Explain how Query-1 and Query-2 will be processed with index file?

 Seek time = movement time of disk head.

For Query-1:

- cost without index:

number of seek = 1

number of block transfer = 11

- cost with index:


number of seek = 1

number of block transfer = 1

For Query-2:

- cost without index:

number of seek = 1

number of block transfer = 12

- cost with index:

number of seek = 1

number of block transfer = 4

 What will happen if the search key is a non-primary column/or secondary index?

 Primary index, secondary index query processing time - important for the exam.

Primary Index:

Order of index file = order of dataset

ordered index →

- Dense (~~one to one~~ one to one)
- Sparse (one to many)
 - ↳ applicable only when the dataset are sorted physically.
 - primary index.
 - means, all sparse index are primary index.

Query Cost:

Sparse > Dense

Query-3:

```
SELECT *  
FROM Instructor  
WHERE ID = 76543
```

 Question-3: find out the cost
using sparse index?

⇒ seek = 1
block transfer = 5

 What will happen when the search key exist on the
sparse index directly?

 What will happen in the case of range access?

✳️ Secondary Indices

- secondary index have to be derive.

✳️ Query-1:

```
SELECT *  
FROM Instructor  
WHERE Salary = 80000
```

Cost:

with index:

seek = 2

Block transfer = 2

without index:

seek = 1

Block transfer = 12

✳️ Query-2:

```
SELECT *  
FROM Instructor  
WHERE Salary >= 65000
```

Cost:

with index:

seek = 9

Block transfer = 9

without index:

seek = 1

Block transfer = 12

✳️ Seek time is almost 10x of Block transfer time.

✳️ B+ tree index file

✳️ Three types of nodes:

- i. Root
- ii. Internal
- iii. Leaf

Number of points = n

Number of maximum search key in a block/node = $n-1$

minimum = $\lceil \frac{n-1}{2} \rceil$

usually it is 100

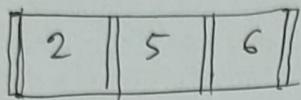
④ $n=4$

Data = 2, 5, 6, 7, 10

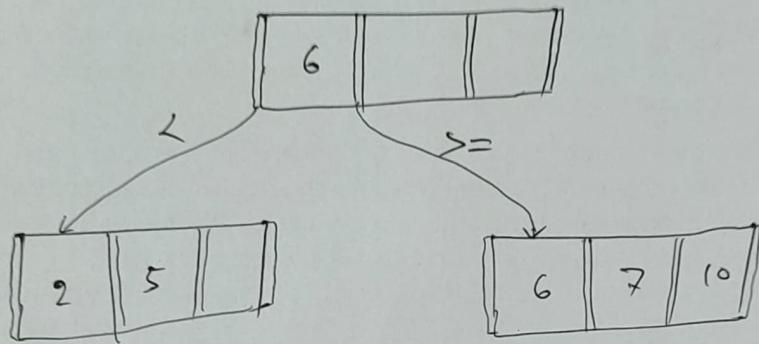
min = 2

max = 3

⇒



⇒



④ Root & internal, there will be no duplicate search-key.

But all the search-key will be available in the leaf node And leaf nodes are in sorted order.

④ Advantage and Disadvantage

Slide - 34

④ Query - 3:

SELECT *

FROM Instructor

WHERE Name > "Crick" ~~and~~

and Name < "Kim"

cost:

seek = 49

Block Transfer = 49

~~Q~~ Query - 4:

```
SELECT *
FROM Instructor
WHERE Name = "Singh"
```

Cost :

seek = 4

Block Transfer = 4

Quiz 1
upto this
30.09.2024

L-06 / 11.09.2024 /

~~Q~~ Query Processing & Optimization

~~Q~~ Algebra:

~~S~~ → Select

Π → Project

~~Q~~ Query - 1:

```
SELECT ID
FROM Student
```

Algebra:

$\Pi_{ID}(\text{Student})$

~~Q~~ Query - 2:

```
SELECT *
FROM Student
WHERE CGPA > 3.0
```

Algebra:

$\sigma_{CGPA > 3.0}(\text{Student})$

~~S~~ ⇒ for all column.

Π ⇒ for some specific column.

Basic Steps

Slide - 3

→ Syntax analyzer

- i. Parse and Translate to relational algebra expression
- ii. Optimize the query with the help of statistic info of database.
- iii. Evaluate.

 If intermediate result is low then the cost will be low.

Query Evaluation:

⇒ many factors contribute to the time cost

- disk access
- CPU
- network communication

But we need to consider only the response time or resource consumption.

 t_T = time to transfer one block

t_s = time for one seek.

 Seek time is usually 40-50 times than block transfer time.



A 1 (Linear Search)

- Scan all the block

$$\Rightarrow t_s + b_n t_T$$

Problem -1: Secondary Index

Size = 2000 blocks

seek = 5 ms

block transfer = 0.1 ms

Query:

SELECT *

FROM student

WHERE ~~CGPA~~ CGPA > 3.5

\Rightarrow cost:

$$b_n = 2000$$

$$\text{cost} = t_s + b_n t_T$$

$$= (5 + 2000 \cdot 0.1) \text{ ms}$$

$$= 205 \text{ ms}$$

Problem -2: Primary Index

Query:

SELECT *

FROM Student

WHERE ID < 1801111042

} Result = 800 blocks

$$\therefore b_n = 800$$

$$\hookrightarrow \text{cost} = t_s + b_n t_T$$

$$= (5 + 800 \cdot 0.1) \text{ ms}$$

$$= 85 \text{ ms}$$

(*) A2 (clustering index equality on key)

- Primary index
 - B^+ tree
 - height h_i given
- retrieve single record.

$$\text{cost} = \underbrace{h_i * (\ell_s + \ell_T)}_{\text{index cost}} + \underbrace{(\ell_s + \ell_T)}_{\text{data copy cost}}$$

Query - 35

```
SELECT *
FROM Student
WHERE ID = 1520303042
```

$$\text{cost} = h_i * (\ell_s + \ell_T) + (\ell_s + \ell_T)$$

(*) A3 (clustering index, equality on nonkey)

- secondary index

- retrieve multiple records, consecutive
- number of block containing matching record.

$$\text{cost} = h_i * (\ell_s + \ell_T) + \ell_s + b * \ell_T$$

Query:

```
SELECT *
FROM Student
WHERE City = "Uttara"
```

Number of student = 400 (uffara)
 each record size = 400 bytes
 each block size = 4000 bytes
 record per block = $4000/400 = 10$
 total block need to transfer = $400/10 = 40$

$$B^+ \text{ tree height} = 2$$

$$\begin{aligned} \text{cost} &= h_i * (\ell_s + \ell_T) + \ell_s + b * \ell_T \\ &= 2 * (10 + 0.2) + 10 + 40 * 0.2 = 38.4 \end{aligned}$$

★ A4 (secondary index, equality on key/non-key) number of record that matched.

$$\text{cost} \Rightarrow (h_i + n) * (t_s + t_r)$$

$$= \underbrace{h_i * (t_s + t_r)}_{\text{index cost}} + \underbrace{n(t_s + t_r)}_{\text{block transfer, data cost}}$$

★ A7 (conjunctive selection using one index)

- apply A4 for each indices and find out the best one.

Problem - 5:

For CGPA

$$\begin{aligned}\text{cost} &= (h_i + n) * (t_s + t_r) \\ &= (2 + 5) * (10 + 0.2) \\ &= 71.4 \text{ ms}\end{aligned}$$

for City

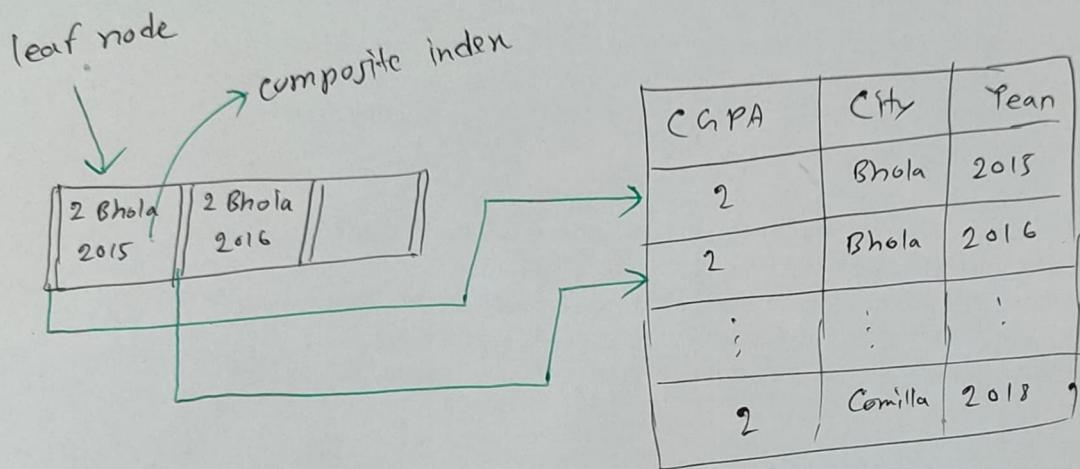
$$\begin{aligned}\text{cost} &= (h_i + n) * (t_s + t_r) \\ &= (2 + 20) * (10 + 0.2) \\ &= 224.4 \text{ ms}\end{aligned}$$

For Year,

$$\begin{aligned}\text{cost} &= (h_i + n) * (t_s + t_r) \\ &= (1 + 4000) * (10 + 0.2) \\ &= 40810.2 \text{ ms}\end{aligned}$$

A8 (conjunctive selection using composite index)

- if composite index available, then it will be faster than others.
- fully optimized, but require more storage.



Query:

SELECT *

FROM Student

WHERE CGPA = 2 AND City = "Bhola" AND Year = 2015

$$\text{cost} = \text{index cost} + \text{data cost}$$

$$= (h_i + n) * (t_s + t_T)$$

A9 (conjunctive selection using intersection of identifiers)

- just use the index attr and all pointers to find out the data location.
- we do not need to transfer data to intersect.
- we can simply intersect the pointers of the buckets of index.

Cost = cost of each index to transfer pointer list
 +
 data cost

Problem -7:

$$\begin{aligned} \text{cost} &= (\sum(h_i) + n) * (t_s + t_T) \\ &= (2+2+1+1) * (10 + 0.2) \\ &= 61.2 \text{ ms.} \end{aligned}$$

1-08 / 23.09.2024 /

④ Recap Algorithm.

④ student (ID, Name, Street, City, CGPA)

Index given = primary index only (ID)
 = secondary index on (Name, Street, City)

④ Query-1:

SELECT *

FROM student

WHERE ID = 1234567

Find the algorithm?

$\Rightarrow A2$

Q1) Query - 2:

```
SELECT *  
FROM Student  
WHERE CGPA = 3.5
```

⇒ A1

- CGPA is not in the index list,
so we need to perform linear
search

Q2) Query - 3:

```
SELECT *  
FROM Student  
WHERE ID > 2345678  
AND Name = "Abdullah"
```

⇒ A7, A9

- composite index not
available, so A8 not
possible.

Q3) Sorting

We may build an index on the relation, and then use the index to read the relation in sorted order. May lead to one disk block access for each tuple.

- we can copy the or transfer the relation to memory and apply quick-sort.
- But if a relation don't fit in the memory, what to do?
- external sort-merge is a good solution for this problem.

~~(A)~~ External Sort-Merge:

N = number of runs

M = memory size in terms of run.

possible to merge = $M-1$ runs *store the merge result.*



Given,

$$N = 90$$

$$M = 10$$

$$\therefore \text{Run required} = \frac{90}{10-1} = \frac{90}{9} = 10 \text{ times}$$

Then, we can merge the 9 runs at once

So total Run time = $9+1 = 10$ times

And the last one will be showed as output to the terminal, no need to store in the disk.

~~(B)~~ Given,

Number of block = 900

Memory = 10 block

$$\Rightarrow \text{Number of Runs} = \left\lceil \frac{900}{10} \right\rceil = 90$$

Total block = b_n

memory size = M

$$\text{Number of initial run} = \left\lceil \frac{b_n}{M} \right\rceil$$

$$1^{\text{st}} \text{ merge, number of runs} = \left\lceil \left\lceil \frac{b_n}{m} \right\rceil / m-1 \right\rceil$$

$$2^{\text{nd}} \text{ merge, number of runs} = \left\lceil \left\lceil \frac{b_n}{m} \right\rceil / m-1 / m-1 \right\rceil$$

$$= \left\lceil \left\lceil \frac{b_n}{m} \right\rceil / (m-1)^2 \right\rceil$$

Therefore,

$$\text{Number of merge pass} = \left\lceil \log_{m-1} \left\lceil \frac{b_n}{m} \right\rceil \right\rceil$$

$$P = m^n$$

$$n = \log_m P$$

L-09/125.09.2029/

$$\begin{array}{lll}
 R_1 & R'_1 & R''_1 \\
 R_2 & R'_2 & R''_2 \\
 R_3 & R'_3 & R''_3 \\
 \vdots & \vdots & \vdots \\
 \vdots & \vdots & \vdots \\
 \vdots & \vdots & \vdots \\
 R & R\left(\frac{N}{m-1}\right) & R\left(\frac{N}{(m-1)^2}\right)
 \end{array}$$

$$\frac{N}{m-1} \quad \frac{\left\lceil \frac{B_n}{m} \right\rceil}{m-1} \quad \frac{\left\lceil \frac{B_n}{m} \right\rceil}{(m-1)^2}$$

$$\textcircled{O} \text{ Number of merge pass} = \log_{m-1} \left\lceil \frac{B_n}{m} \right\rceil$$

$$\textcircled{O} \text{ Total Blocks} = B_n$$

⊗ Block transfer for run creation = $2 \cdot B_n$

⊗ Total Block Transfer = $2 B_n + 2 B_n * \log_{(m-1)} \lceil \frac{B_n}{m} \rceil$

Run creation Merge Pass

$- B_n$
Last result will not store to the disk.

∴ Total Seek = Run Creation + Merge Pass

$$= 2 \lceil \frac{B_n}{m} \rceil + 2 B_n \lceil \log_{m-1} \frac{B_n}{m} \rceil - B_n$$

⊗ Join Operation

- Nested Loop
- Block nested loop

⊗ student (ID, Name, Street, City, CGPA)

Takes (ID, CourseID, Semester, Year, Grade)

⇒ Find students and corresponding course information.

∴ SQL -

```
SELECT *
FROM student INNER JOIN Takes ON student.ID =
Takes.ID
```

Algebra: Student \bowtie Takes \Rightarrow What will be the cost

\Rightarrow Find out cost of Takes \bowtie Student?

Will they takes different cost? why?

\Rightarrow Estimated block transfer

= number of tuples in first relation * number of block in second relation + number of block in first relation

$$= n_r \times b_s + b_n$$

$n \bowtie s$

\Rightarrow Number of seek

= number of tuples in first relation

+ number of blocks in first relation

$$= n_r + b_n$$

Best Case
B.T. = $b_s + b_n$
Seek = 2

Therefore,

~~for~~ ✓ Student \bowtie Takes

$$\begin{aligned} \text{block transfer} &= 20,000 \times 400 + 2,000 \\ &= 8,02,000 \end{aligned}$$

$$\text{seek} = 20,000 + 2,000$$

$$= 22,000$$

Takes \bowtie Student

$$\text{Block Transfer} = 40,000 \times 2,000 + 400 = 8,00,00,400$$

$$\text{seek} = 40,000 + 400 = 40,400$$