

Quiz-1

⊗ Block nested loop

⊗ Worst case!

Block Transfer = number of block of first relation \times
number of block of second relation
+ number of block of first relation

$$= b_r \times b_s + b_r$$

$$\text{Seek} = b_r + b_s$$

Best Case!

$$\text{Block Transfer} = b_r + b_s$$

$$\text{Seek} = 2$$

⊗ Evaluation of Expression

i. Materialization!

- generate results of an expression whose inputs are relations or already computed.
- Materialize : store on disk.

ii. Pipelining:

- pass ~~one~~ on tuples to parent operations even as an operation is being executed.

⊗ Query-1:

```
SELECT Name
FROM Instructor INNER JOIN Department
ON Instructor.dept-name = Department.dept-name
WHERE Building = "Watson"
```

⇒ Algebra:

$$\pi_{\text{Name}} \left(\left(\sigma_{\text{Building} = \text{Watson}} (\text{Department}) \right) \bowtie \text{Instructor} \right)$$

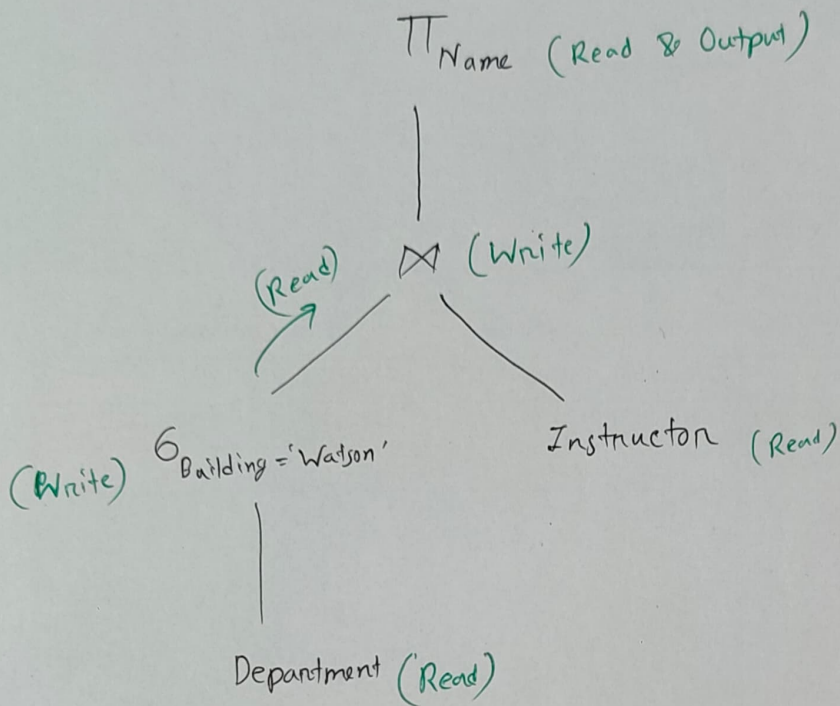
⇒ Execution will start from the inner parenthesis.

But, we will construct the execution tree from outer parenthesis.

As in our classical computer, we call function in another function, like nested, then we get the result on return call one by one.

Execution tree is like that, it first goes to the leaf, & once leaf is reached, it starts to return the result to parent and finally we get the output from the root.

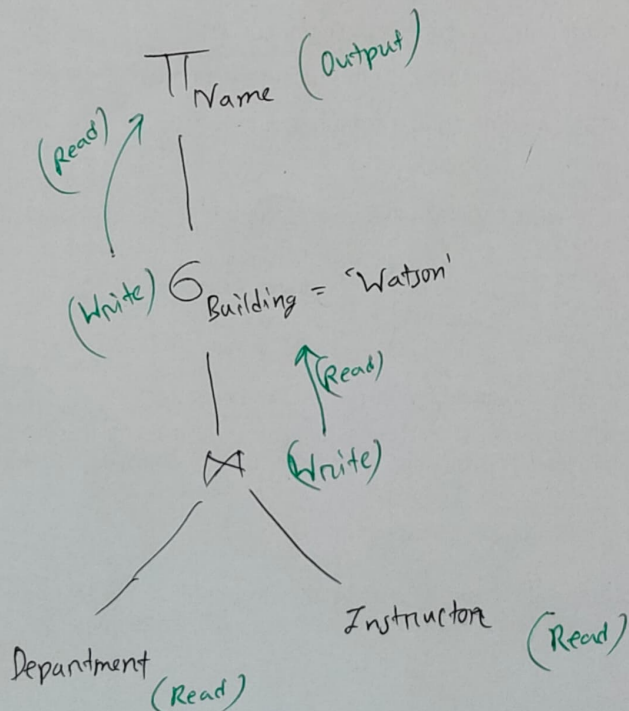
Execution tree:



Algebra:

$$\Pi_{Name} \left(\sigma_{Building = 'Watson'} (Department \bowtie Instructor) \right)$$

Execution Tree:



L-11 / 02.10.2024 /

Midterm

23.10.2024

upto Transaction

Lecture Slide - 0,1,2

⊗ From the sample Question:


Ⓐ

Algebra:

$\pi_{sid, name} (\sigma_{city = 'Dhaka' \wedge year = 2023} (Student \bowtie Take))$

$\pi_{sid, name} (Output)$
|
(Read)

$\sigma_{city = Dhaka \wedge year = 2023} (Write)$



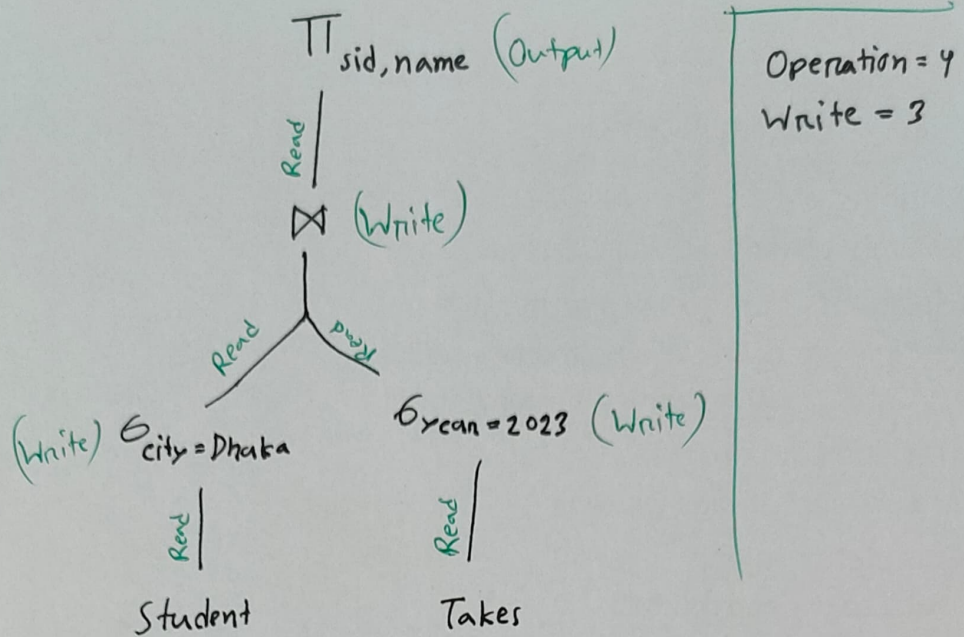
Operation = 3

Write = 2

still it will take more
time, because of large
join operation.

⇒ Optimized:

$\pi_{sid, name} (\sigma_{city = Dhaka} (Student) \bowtie \sigma_{year = 2023} (Take))$



⊗ For read & write, cost is high for materialization.

$$\text{Overall Cost} = \sum \text{operation cost} + \sum \text{write cost}$$

⊗ Problem-13:

$$\text{Overall Cost} = \sum \text{operation cost} + \sum \text{write cost}$$

$$= (1S + 20B + 10S + 100B + 1S + 50B) \\ + (1S + 10B + 1S + 50B)$$

$$= 14S + 230B$$

⊗ Materialization method is ^Aso costly for read/write.
There is another method to reduce this read write waiting time.

Double buffering:

Double buffering:

- two buffer for each operation
- when 1 buffer is full, it start to write on disk and in the mean time other buffer store the operation result.

Pipelining

- ⇒ No result will be store on disk. Once one tuple found that match with the condition, it will send the data to the next operation directly.
- ⇒ Pipelining is not possible for sorting or hash-join.

⊗ Pipelining method can be executed in two ways:

- ⊗ Pull model ← (i) Demand driven ^{lazy evaluation} ⇒ Bottom level wait for the request from Top level.
- ⊗ Push model ← (ii) Producer driven ^{eager} ⇒ Bottom level generate output and push it to top level for further operation.
 - maintain buffer - puts & remove.

Transaction Manager

Transaction:

- is a unit of program execution that accesses and possibly updates various data items.

Maintain:

- from hardware failure, system crash
- concurrent execution of multiple transaction.

ACID Properties:

Atomicity:

- the system should ensure that, updates of a partially executed transaction are not reflected in the database.

Consistency:

- when the transaction completes successfully the database must be consistent.
- Like: sum of balance of all account must be same as before.

Isolation:

- when multiple transaction are executed concurrently, each transaction must be unaware of other transaction.
- one transaction can't effect others.

Durability:

- after a transaction completes successfully, the changes it has made to the database persist, even if there are system failures.

L-12 / 07.10.2024 /

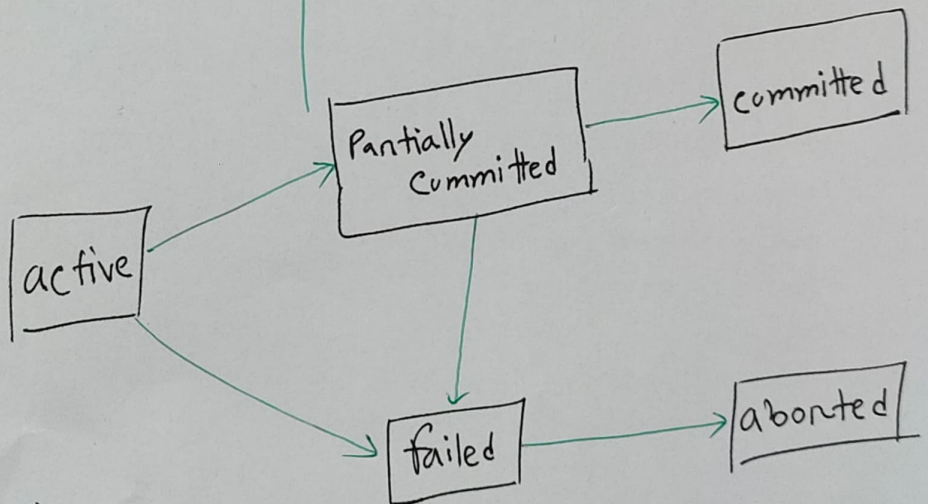
⊗ Transaction States

- Active
- Partially Committed
- Failed
- Aborted → Restart kill
- Committed


T:

START
READ(A)
 $A = A - 10,000$
WRITE(A)
COMMIT → Partially Committed

} Failed State



⊗ Serial Execution:




⊗ Throughput high/low?
⇒ Low.

⊗ No integrity problem.

⇒ Integrity = Yes

⊗ Throughput ⇒ number of transaction per second.

* Concurrent Transaction:



- integrity problem may exist.
- maintained by "concurrency control schemes" - mechanism to achieve isolation.

* Throughput high/low?
⇒ High.

* Throughput also means resource utilization. Where resource utilization is high/man/100%, then throughput must be high.

* Advantage of concurrency:

- increased processor & disk utilization.
- reduces average response time.

* Schedule $\left\{ \begin{array}{l} \rightarrow \text{Serial} \\ \rightarrow \text{concurrent} \end{array} \right.$

- transaction sets must be consistent.
- preserve the order in which the transaction appear.

*

- Schedule-1 ⇒ serial ⇒ database consistent.
- schedule-2 ⇒ serial ⇒ database consistent
- schedule-3 ⇒ concurrent ⇒ consistent
- schedule-4 ⇒ concurrent ⇒ not consistent.

⊗ Serializability! $\begin{cases} \rightarrow \text{conflict} \\ \rightarrow \text{view} \end{cases}$

- concurrent schedule can be serializable, if the schedule is equivalent to a serial schedule.


\Rightarrow

consistency exist \Rightarrow serializable

" not exist \Rightarrow not serializable.

⊗ Let, there are three transaction: T_1, T_2, T_3

concurrent schedule:



Serial transaction possibility:

T_1	T_2	T_3
T_1	T_3	T_2
T_2	T_1	T_3
T_2	T_3	T_1
T_3	T_1	T_2
T_3	T_2	T_1

\Rightarrow if concurrent schedule is equivalent to any of the serial schedule, then it is serializable.

\Rightarrow And if it is serializable, then it must be consistent.

⊗ Schedule only depends on read/write operation.

⊗ Conflicting instruction:

- occur when read/write operation requested in the same object/account.

- ⊛ Between two transaction on the same account:
- two or more than.
 - if both are read operation, then no conflict.
 - in all other combination, there must be a conflict.
 - write operation must be done independently.

Example! slide - 20, 21

L-13 / 09.10.2024 /

⊛ Example:

T_1	T_2
READ(P)	
READ(Q)	WRZTE(Q)
	WRZTE(R)
WRITE(Q)	
WRZTE(R)	

⇒

T_1	T_2
READ(P)	WRZTE(Q)
READ(Q)	WRZTE(R)
WRZTE(Q)	
WRZTE(R)	

⇒ Serializable

⊛

T_1	T_2	T_3
READ(A)		
READ(B)	READ(A)	
	READ(B)	
		WRITE(D)

⇒ Possible to serialize in T_1, T_2, T_3

⇒ Serializable

all read operation
- no conflict

→ write operation
but no other operation
on D.



T_1	T_2	T_3
READ(P)		
	WRITE(Q)	
READ(Q)	WRITE(R)	
WRITE(Q)		
		READ(R)
WRITE(R)		

$\Rightarrow T_2, T_3, T_1$

= Serializable

* Swapping method takes too much time and space.


There is another algorithm regarding the graph theory.

\Rightarrow Precedence Graph

- vertices are the transaction


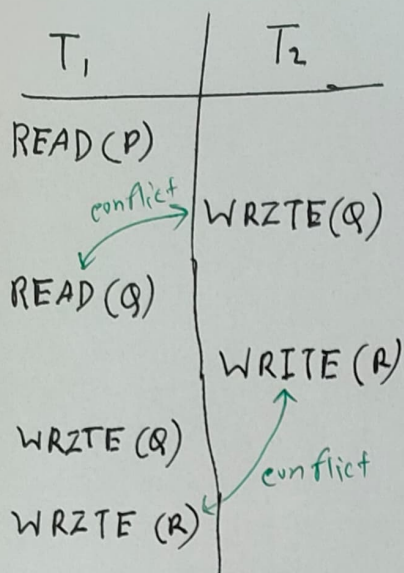


T_3	T_4
READ(Q)	
	WRITE(Q)
WRITE(Q)	



T_1	T_2	T_3
READ(A)		
	READ(A)	
READ(B)		
	READ(B)	
		WRITE(D)

\Rightarrow No conflict operation
- no graph
- serializable.



Same areh for two conflict because of T_2 is in up than T_1

⇒ No circle
- serializable

⊗ If precedence graph is acyclic, then serializability order can be obtained by a topological sorting of the graph.

Midterm Syllabus
upto this

⊗ Concurrency Control

⊗ Lock mechanism to control concurrent access:

two modes of lock:

(i) exclusive (X) ⇒ Read + Write permission

(ii) shared (S) ⇒ Only read permission

For concurrent lock requests

	S	X
S	true	false
X	false	false

if both are 'S' mode, then allowed other-wise not allowed.